

2020 Summer School on Effective HPC for Climate and Weather

Input/Output and Middleware

Luciana Pedro, Julian Kunkel

Department of Computer Science, University of Reading

18 June 2020



- 1 NetCDF Files and C
- 2 NetCDF Utilities
- 3 NetCDF Operators
- 4 Parallel I/O
- 5 Practising
- 6 Starting with NetCDF

Disclaimer: This material reflects only the author's view and the EU-Commission is not responsible for any use that may be made of the information it contains

Learning Objectives

- Execute programs in C that read and write NetCDF files in a metadata-aware manner
- Analyze, manipulate and visualise NetCDF data
- Implement an application that utilizes parallel I/O to store and analyze data

References

- The files and data used in this presentation were collected on the Unidata website.
 - ▶ <https://www.unidata.ucar.edu/>
- All files used here are available in the following Git Repository:
 - ▶ <https://github.com/ESiWACE/io-training>
- These files are also available with the NetCDF main installation, in the directory `examples`.
- For more information about how to install NetCDF in your personal computer, from scratch, check Section 6.


File Reference: `simple_xy_wr.c`

- This is an example program demonstrating a simple 2D write. It is intended to illustrate the use of the netCDF C API.

▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_wr_8c.html

▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_wr_8c_source.html

- Dependency graph for `simple_xy_wr`:



File simple_xy_wr.c: Header and Constants Declaration

```
#include <stdlib.h>
#include <stdio.h>
#include <netcdf.h>

/* This is the name of the data file we will create. */
#define FILE_NAME "simple_xy.nc"

/* We are writing 2D data, a 6 x 12 grid. */
#define NDIMS 2
#define NX 6
#define NY 12

/* Handle errors by printing an error message and exiting with a
 * non-zero status. */
#define ERRCODE 2
#define ERR(e) {printf("Error: %s\n", nc_strerror(e)); exit(ERRCODE);}

int main()
{
    ...
    ...
    ...
}
```

File simple_xy_wr.c: Variables Declaration

```

...
...
...

int main()
{
    /* When we create netCDF variables and dimensions, we get back an
     * ID for each one. */
    int ncid, x_dimid, y_dimid, varid;
    int dimids[NDIMS];

    /* This is the data array we will write. It will be filled with a
     * progression of numbers for this example. */
    int data_out[NX][NY];

    /* Loop indexes, and error handling. */
    int x, y, retval;

    ...
    ...
    ...
}

```

File simple_xy_wr.c: Creating (loading!) data

```
...

int main()
{
    ...

    /* Create some pretend data. If this wasn't an example program, we
     * would have some real data to write, for example, model
     * output. */
    for (x = 0; x < NX; x++)
        for (y = 0; y < NY; y++)
            data_out[x][y] = x * NY + y;

    ...
}
```


File simple_xy_wr.c: Creating the NetCDF file

```

...

int main()
{
    ...

    /* Always check the return code of every netCDF function call. In
     * this example program, any retval which is not equal to NC_NOERR
     * (0) will cause the program to print an error message and exit
     * with a non-zero return code. */

    /* Create the file. The NC_CLOBBER parameter tells netCDF to
     * overwrite this file, if it already exists.*/
    if ((retval = nc_create(FILE_NAME, NC_CLOBBER, &ncid)))
        ERR(retval);

    ...
}

```

File simple_xy_wr.c: Defining the dimensions

```

...

int main()
{
    ...

    /* Define the dimensions. NetCDF will hand back an ID for each. */
    if ((retval = nc_def_dim(ncid, "x", NX, &x_dimid)))
        ERR(retval);
    if ((retval = nc_def_dim(ncid, "y", NY, &y_dimid)))
        ERR(retval);

    /* The dimids array is used to pass the IDs of the dimensions of
       * the variable. */
    dimids[0] = x_dimid;
    dimids[1] = y_dimid;

    ...
}

```

File simple_xy_wr.c: Defining the variable

```

...

int main()
{
    ...

    /* Define the variable. The type of the variable in this case is
     * NC_INT (4-byte integer). */
    if ((retval = nc_def_var(ncid, "data", NC_INT, NDIMS,
                            dimids, &varid)))
        ERR(retval);

    /* End define mode. This tells netCDF we are done defining
     * metadata. */
    if ((retval = nc_enddef(ncid)))
        ERR(retval);

    ...
}

```

File simple_xy_wr.c: Writing data to the file

```

...

int main()
{
    ...

    /* Write the pretend data to the file. Although netCDF supports
     * reading and writing subsets of data, in this case we write all
     * the data in one operation. */
    if ((retval = nc_put_var_int(ncid, varid, &data_out[0][0])))
        ERR(retval);

    /* Close the file. This frees up any internal netCDF resources
     * associated with the file, and flushes any buffers. */
    if ((retval = nc_close(ncid)))
        ERR(retval);

    ...
}

```

File simple_xy_wr.c: Getting SUCCESS!

```
...  
  
int main()  
{  
    ...  
  
    printf("*** SUCCESS writing example file simple_xy.nc!\n");  
    return 0;  
}
```

Compiling and running the file `simple_xy_wr.c`

■ Create (copy!) and compile the file `simple_xy_wr.c`.

▶ `gcc -I/home/username/local/include simple_xy_wr.c -o simple_xy_wr -L/home/username/local/lib -lnetcdf`

▶ LR: What does that mean?!

■ Run the file `simple_xy_wr`.

▶ `./simple_xy_wr`

▶ `*** SUCCESS writing example file simple_xy.nc!`

■ Check that the file `simple_xy.nc` is in your directory.


File Reference: `simple_xy_rd.c`

- This is a simple example which reads a small dummy array that was written by `simple_xy_wr.c`.

▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_rd_8c.html

▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_rd_8c_source.html

- Dependency graph for `simple_xy_rd`:



File simple_xy_rd.c

```
int main()
{
    /* Open the file. NC_NOWRITE tells netCDF we want read-only access
     * to the file.*/
    if ((retval = nc_open(FILE_NAME, NC_NOWRITE, &ncid)))
        ERR(retval);

    /* Get the varid of the data variable, based on its name. */
    if ((retval = nc_inq_varid(ncid, "data", &varid)))
        ERR(retval);

    /* Read the data. */
    if ((retval = nc_get_var_int(ncid, varid, &data_in[0][0])))
        ERR(retval);

    /* Check the data. */
    for (x = 0; x < NX; x++)
        for (y = 0; y < NY; y++)
            if (data_in[x][y] != x * NY + y)
                return ERRCODE;

    /* Close the file, freeing all resources. */
    if ((retval = nc_close(ncid)))
        ERR(retval);
}
```


Reading the file `simple_xy.nc`

■ Check that the file `simple_xy.nc` is in your directory.

■ Create (copy!), compile and run the file `simple_xy_rd.c`.

```
▶ gcc -I/home/username/local/include simple_xy_rd.c -o simple_xy_rd -L/home/username/local/lib -lnetcdf
```


■ Run the file `simple_xy_rd`.

```
▶ ./simple_xy_rd
```

```
▶ *** SUCCESS reading example file simple_xy.nc!
```

ncdump and ncgen

- ncdump and ncgen are inverses:



- Used together, ncdump and ncgen can accomplish simple netCDF manipulations with little or no programming.

Editing a NetCDF File


- To edit metadata or data in a netCDF file:



- ▶ Use `ncdump` to convert netCDF file to CDL.
- ▶ Use a text editor to make desired change to CDL.
- ▶ Use `ncgen` to turn modified CDL back into netCDF file.
- ▶ **Note:** This option is not practical for huge netCDF files or if one intend to modify lots of files. For that, need to write a program using netCDF library.

Creating a NetCDF File

- To create a new netCDF file with lots of metadata:



- ▶ Use a text editor to write a CDL file with lots of metadata but little or no data.
- ▶ Use `ncgen` to generate corresponding C or Fortran program for writing netCDF.
- ▶ Insert appropriate netCDF **var_put** calls for writing data.
- ▶ Compile and run program to create netCDF file.
- ▶ Use `ncdump` to verify result.

Using ncdump

■ Inspect the file simple_xy.nc using ncdump

► `ncdump simple_xy.nc`

► LR: Only works like that in my laptop:

► `/home/lucy/netcdf/netcdf-c-4.7.4/ncdump/ncdump simple_xy.nc`

► LR: It should be just `ncdump file`, but `esdm` is on my way and I don't know how to make another link.

NetCDF CDL Format

```
netcdf simple_xy {  
  dimensions:  
    x = 6 ;  
    y = 12 ;  
  variables:  
    int data(x, y) ;  
  data:  
  
    data =  
      0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11,  
      12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23,  
      24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35,  
      36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47,  
      48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59,  
      60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71 ;  
}
```

Using ncgen

■ Create a NetCDF file using ncgen and the CDL output



```
/home/lucy/netcdf/netcdf-c-4.7.4/ncdump/ncdump simple_xy.nc > simple_xy_t
```



```
more simple_xy_test.cdl
```



```
/home/lucy/netcdf/netcdf-c-4.7.4/ncgen/ncgen -b simple_xy_test.cdl
```



```
cmp simple_xy_test.nc simple_xy.nc
```



LR: Fix [ncdump](#)!

Creating the C File

■ Create a C file using ncgen and the CDL output



```
/home/lucy/netcdf/netcdf-c-4.7.4/ncgen/ncgen -lc simple_xy_test.cdl > simple_xy_test.c
```



```
more simple_xy_test.c
```



What is the difference between the files `simple_xy_test.c` and `simple_xy_wr.c`?



LR: Fix ncgen!



```
cmp simple_xy_test.c simple_xy_wr.c
```



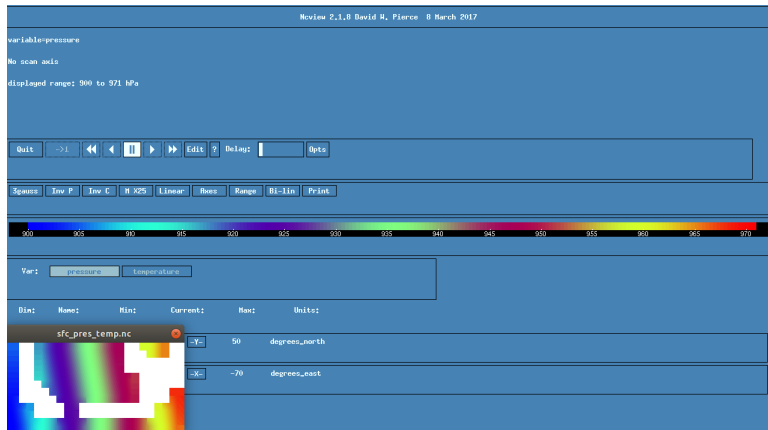
```
meld simple_xy_test.c simple_xy_wr.c
```


Starting all over again!

- `gcc -I/home/username/local/include simple_xy_test.c
-o simple_xy_test -L/home/username/local/lib -lnetcdf`
- `mv simple_xy_test.nc simple_xy_test2.nc`
- `./simple_xy_test`
- `cmp simple_xy_test.nc simple_xy_test2.nc`

ncview

- LR: Installation is fine.
- LR: What does it do?



NetCDF Operators (NCO)LR: Is there time to explore?

- NCO is a package of command line operators that manipulates generic netCDF data and supports some CF conventions. The NCO utilities are:

ncap2	arithmetic processor
ncatted	attribute editor
ncbo	binary operator
ncdiff	differencer
ncea	ensemble averager
ncecat	ensemble concatenator
ncflint	file interpolator

ncks	kitchen sink (extract, cut, paste, print data)
ncpdq	permute dimensions quickly
ncra	running averager
ncrcat	record concatenator
ncrename	renamer
ncwa	weighted averager

- NCO utilities have as a goal being as generic as possible, imposing no limitations on data dimensionality, size, or type. All established national and international climate modelling centers now install and maintain NCO for their users for data post-processing, hyper-slabbing, and serving.

Parallel I/O

Implement an application that utilises parallel I/O to store and analyse data

Files for Practising

■ File `simple_xy_nc4`

- ▶ Write/Read the `simple_xy` file with some of the features of netCDF-4.
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_nc4_wr_8c.html
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_nc4_rd_8c.html

■ File `simple_nc4`

- ▶ Write/Read a file demonstrating some of the features of netCDF-4.
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_nc4_wr_8c.html
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_nc4_rd_8c.html

■ File `sfc_pres_temp_wr`

- ▶ This is an example program which writes/reads surface pressure and temperatures.
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/sfc_pres_temp_wr_8c.html
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/sfc_pres_temp_rd_8c.html

■ File `pres_temp_4D_wr`

- ▶ This is an example program which writes/reads 4D pressure and temperatures.
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/pres_temp_4D_wr_8c.html
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/pres_temp_4D_rd_8c.html

Summary of Actions

- Inspect the read and write files in C code.
- Compile and run the write/read C files.
- Inspect the output NetCDF file (.nc) using `ncdump`.
- Create a CDL file for the NetCDF file.
- Recreate the NetCDF file using `ncgen` and the CDL file.
- Recreate the C file using `ncgen` and the CDL file.
- Visualize the data in the NetCDF file with `ncview`.

Building NetCDF from Scratch

- The usual way of building netCDF requires the HDF5, zlib, and curl libraries.
- Files for the libraries can be found in:

`ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-4`

Installing curl

■ `apt-get install libcurl4-openssl-dev`

Installing zlib

- `wget`
`ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-4/zlib-1.2.8.tar.gz`
 - ▶ Newest version to later use `ncview`
 - ▶ `wget https://sourceforge.net/projects/libpng/files/zlib/1.2.9/zlib-1.2.9.tar.gz`
- `tar -xvzf zlib-1.2.8.tar.gz`
- `cd zlib-1.2.8`
- `mkdir /home/username/local/`
- `./configure --prefix=/home/username/local/`
- `make check install`

Installing HDF5

- `wget`
`ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-4/hdf5-1.8.13.tar.gz`
- `tar -xvzf hdf5-1.8.13.tar.gz`
- `cd hdf5-1.8.13`
- `./configure --with-zlib=/home/username/local/ --prefix=/home/username/local/`
- `make`
- `make check`
- `make install`
 - ▶ `make check install`
 - ▶ If not done separately, it might not work!

Installing NetCDF

- Check the latest version at
<https://www.unidata.ucar.edu/downloads/netcdf/>
- `wget ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-c-4.7.4.tar.gz`
- `tar -xvzf netcdf-c-4.7.4.tar.gz`
- `cd netcdf-c-4.7.4`
- `CPPFLAGS=-I/home/username/local/include`
`LDLFLAGS=-L/home/username/local/lib ./configure`
`--prefix=/home/username/local`
- `make check install`

Finishing the Set Up

■ Link the NetCDF library

- ▶ `export LD_LIBRARY_PATH=/home/username/local/lib/`
- ▶ `sudo ldconfig`

■ Create a new directory (for instance, `/home/username/example`) and create the file from the given source using an editor of your choice.


File Reference: `simple_xy_wr.c`

This is an example program demonstrating a simple 2D write. It is intended to illustrate the use of the netCDF C API.

■ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_wr_8c.html

■ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_wr_8c_source.html

Dependency graph for `simple_xy_wr`:



Compiling and running the file `simple_xy_wr.c`

- Create (copy!) and compile the file `simple_xy_wr.c`.

- ▶ `gcc -I/home/username/local/include simple_xy_wr.c -o simple_xy_wr -L/home/username/local/lib -lnetcdf`

- Run the file `simple_xy_wr`.

- ▶ `./simple_xy_wr`

- ▶ `*** SUCCESS writing example file simple_xy.nc!`

- Check that the file `cmp test.nc simple_xy.nc` is in your directory.

Using ncdump

Inspect the output file `simple_xy.nc` using `ncdump`

- `ncdump simple_xy.nc`
- LR: Only works like that in my laptop:
- `/home/lucy/netcdf/netcdf-c-4.7.4/ncdump/ncdump simple_xy.nc`
- LR: It should be just `ncdump file`, but `esdm` is on my way and I don't know how to make another link

NetCDF CDL Format

```
netcdf simple_xy {
dimensions:
x = 6 ;
y = 12 ;
variables:
int data(x, y) ;
data:

data =
  0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11,
  12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23,
  24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35,
  36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47,
  48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59,
  60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71 ;
}
```


Using ncgen

■ Create a NetCDF file using ncgen and the CDL output



```
/home/lucy/netcdf/netcdf-c-4.7.4/ncdump/ncdump simple_xy.nc > test.cdl
```



```
more test.cdl
```



```
/home/lucy/netcdf/netcdf-c-4.7.4/ncgen/ncgen -b test.cdl
```



```
cmp test.nc simple_xy.nc
```



```
LR: Fix ncdump!
```

Creating the C File

■ Create a C file using ncgen and the CDL output



```
/home/lucy/netcdf/netcdf-c-4.7.4/ncgen/ncgen -lc simple_xy_test.cdl > sim
```



```
more simple_xy_test.c
```



```
cmp simple_xy_test.c simple_xy_wr.c
```



LR: Fix ncgen!

Starting all over again!



```
gcc -I/home/username/local/include simple_xy_test.c -o simple_xy_test -L
```



```
mv simple_xy_test.nc simple_xy_test2.nc
```



```
./simple_xy_test
```



```
cmp simple_xy_test.nc simple_xy_test2.nc
```


File Reference: `simple_xy_rd.c`

This is a simple example which reads a small dummy array, which was written by `simple_xy_wr.c`. It is intended to illustrate the use of the netCDF C API.

■ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_rd_8c.html

■ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_rd_8c_source.html

Dependency graph for `simple_xy_wr`:



Reading the file `simple_xy.nc`

■ Check that the file `simple_xy.nc` is in your directory.

■ Create (copy!), compile and run the file `simple_xy_rd.c`

```
▶ gcc -I/home/username/local/include simple_xy_rd.c -o simple_xy_rd -L/home/username/local/lib -lnetcdf
```


■ Run the file `simple_xy_rd`

```
▶ ./simple_xy_rd
```

```
▶ *** SUCCESS reading example file simple_xy.nc!
```

ncview

- Installation is fine.
- What does it do?



The ESiWACE1/2 projects have received funding from the European Union's Horizon 2020 research and innovation programme under grant agreement No **675191** and No **823988**



Disclaimer: This material reflects only the author's view and the EU-Commission is not responsible for any use that may be made of the information it contains