

Summer School on Effective HPC for Climate and Weather

Input/Output and Middleware

Luciana Pedro, Julian Kunkel

Department of Computer Science, University of Reading

18 June 2020



- 1 Building NetCDF
- 2 File `simple_xy_wr.c`
- 3 File `simple_xy_nc4_wr.c`
- 4 File `simple_nc4_wr.c`
- 5 File `sfc_pres_temp_wr.c`
- 6 File `pres_temp_4D_wr.c`

Disclaimer: This material reflects only the author's view and the EU-Commission is not responsible for any use that may be made of the information it contains

Learning Objectives

- Describe the role of middleware and file formats (Section Middleware)
- Discuss challenges for data-driven research (Section Data)
- Identify typical I/O performance issues and their causes (Section Performance)
- Apply performance models to assess and optimize the application I/O performance (Section Model)
- Design a data model for NetCDF/CF (Section NetCDF)
- Execute programs in C and Python that read and write NetCDF files in a metadata-aware manner (Section Progs)
- Analyze, manipulate and visualise NetCDF data (Section NetCDF2)
- Implement an application that utilizes parallel I/O to store and analyze data (Section Parallel I/O)
- Describe ongoing research activities in high-performance storage (Section Storage)

Building NetCDF from Scratch

- The usual way of building netCDF requires the HDF5, zlib, and curl libraries.
- Files for the libraries can be found in:

`ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-4`

Installing curl

■ apt-get install libcurl4-openssl-dev

Installing zlib

- wget
ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-4/zlib-1.2.8.tar.gz
 - ▶ Newest version to later use ncview
 - ▶ wget https://sourceforge.net/projects/libpng/files/zlib/1.2.9/zlib-1.2.9.tar.gz
- tar -xvzf zlib-1.2.8.tar.gz
- cd zlib-1.2.8
- mkdir /home/username/local/
- ./configure --prefix=/home/username/local/
- make check install

Installing HDF5

- `wget`
`ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-4/hdf5-1.8.13.tar.gz`
- `tar -xvzf hdf5-1.8.13.tar.gz`
- `cd hdf5-1.8.13`
- `./configure --with-zlib=/home/username/local/ --prefix=/home/username/local/`
- `make`
- `make check`
- `make install`
 - ▶ `make check install`
 - ▶ If not done separately, it might not work!

Installing NetCDF

- Check the latest version at
`https://www.unidata.ucar.edu/downloads/netcdf/`
- `wget ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-c-4.7.4.tar.gz`
- `tar -xvzf netcdf-c-4.7.4.tar.gz`
- `cd netcdf-c-4.7.4`
- `CPPFLAGS=-I/home/username/local/include`
`LDFLAGS=-L/home/username/local/lib ./configure`
`-prefix=/home/username/local`
- `make check install`

Finishing the set up

■ Link the NetCDF library

- ▶ `export LD_LIBRARY_PATH=/home/username/local/lib/`
- ▶ `sudo ldconfig`


■ Create a new directory (for instance, /home/username/example) and create the file from the given source using an editor of your choice.

File Reference: simple_xy_wr.c

This is an example program demonstrating a simple 2D write. It is intended to illustrate the use of the netCDF C API.

- https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_wr_8c.html
- https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_wr_8c_source.html

Dependency graph for simple_xy_wr:



File simple_xy_wr.c: Header and Constants Declaration

```
#include <stdlib.h>
#include <stdio.h>
#include <netcdf.h>

/* This is the name of the data file we will create. */
#define FILE_NAME "simple_xy.nc"

/* We are writing 2D data, a 6 x 12 grid. */
#define NDIMS 2
#define NX 6
#define NY 12

/* Handle errors by printing an error message and exiting with a
 * non-zero status. */
#define ERRCODE 2
#define ERR(e) {printf("Error: %s\n", nc_strerror(e)); exit(ERRCODE);}

int
main()
{
    ...
    ...
    ...
}
```

File simple_xy_wr.c: Variables Declaration

```
...
...
...

int
main()
{
    /* When we create netCDF variables and dimensions, we get back an
     * ID for each one. */
    int ncid, x_dimid, y_dimid, varid;
    int dimids[NDIMS];

    /* This is the data array we will write. It will be filled with a
     * progression of numbers for this example. */
    int data_out[NX][NY];

    /* Loop indexes, and error handling. */
    int x, y, retval;

    ...
    ...
    ...
}
```

File simple_xy_wr.c: Creating (loading!) data

```
...

int
main()
{
    ...

    /* Create some pretend data. If this wasn't an example program, we
     * would have some real data to write, for example, model
     * output. */
    for (x = 0; x < NX; x++)
        for (y = 0; y < NY; y++)
            data_out[x][y] = x * NY + y;

    ...
}
```

File simple_xy_wr.c: Creating the NetCDF file

```
...

int
main()
{
    ...

    /* Always check the return code of every netCDF function call. In
     * this example program, any retval which is not equal to NC_NOERR
     * (0) will cause the program to print an error message and exit
     * with a non-zero return code. */

    /* Create the file. The NC_CLOBBER parameter tells netCDF to
     * overwrite this file, if it already exists.*/
    if ((retval = nc_create(FILE_NAME, NC_CLOBBER, &ncid)))
        ERR(retval);

    ...
}
```

File simple_xy_wr.c: Defining the dimensions

```
...

int
main()
{
    ...

    /* Define the dimensions. NetCDF will hand back an ID for each. */
    if ((retval = nc_def_dim(ncid, "x", NX, &x_dimid)))
        ERR(retval);
    if ((retval = nc_def_dim(ncid, "y", NY, &y_dimid)))
        ERR(retval);

    /* The dimids array is used to pass the IDs of the dimensions of
       * the variable. */
    dimids[0] = x_dimid;
    dimids[1] = y_dimid;

    ...
}
```

File simple_xy_wr.c: Defining the variable

```
...

int
main()
{
    ...

    /* Define the variable. The type of the variable in this case is
     * NC_INT (4-byte integer). */
    if ((retval = nc_def_var(ncid, "data", NC_INT, NDIMS,
                           dimids, &varid)))
        ERR(retval);

    /* End define mode. This tells netCDF we are done defining
     * metadata. */
    if ((retval = nc_enddef(ncid)))
        ERR(retval);

    ...
}
```


File simple_xy_wr.c: Writing data to the file

```
...

int
main()
{
    ...

    /* Write the pretend data to the file. Although netCDF supports
     * reading and writing subsets of data, in this case we write all
     * the data in one operation. */
    if ((retval = nc_put_var_int(ncid, varid, &data_out[0][0])))
        ERR(retval);

    /* Close the file. This frees up any internal netCDF resources
     * associated with the file, and flushes any buffers. */
    if ((retval = nc_close(ncid)))
        ERR(retval);

    ...
}
```

File simple_xy_wr.c: Getting SUCCESS!

```
...

int
main()
{
    ...

    printf("*** SUCCESS writing example file simple_xy.nc!\n");
    return 0;
}
```

Compiling and running the file simple_xy_wr.c

Create (copy!) and compile the file simple_xy_wr.c

■ `gcc -I/home/username/local/include simple_xy_wr.c -o simple_xy_wr -L/home/username/local/lib -lnetcdf`

► What does that mean?!

■ Run the file simple_xy_wr

► `./simple_xy_wr`

► `*** SUCCESS writing example file simple_xy.nc!`

■ Check that the file `cmp test.nc simple_xy.nc` is in your directory

► `ls`

Using ncdump

Inspect the output file `simple_xy.nc` using `ncdump`

- `ncdump simple_xy.nc`
- Only works like that in my laptop:
- `/home/lucy/netcdf/netcdf-c-4.7.4/ncdump/ncdump simple_xy.nc`
- (It should be just `ncdump file`, but `esdm` is on my way and I don't know how to make another link)

NetCDF CDL Format

```
netcdf simple_xy {
dimensions:
x = 6 ;
y = 12 ;
variables:
int data(x, y) ;
data:

data =
  0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11,
  12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23,
  24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35,
  36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47,
  48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59,
  60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71 ;
}
```

Using ncgen

Create a NetCDF file using ncgen and the CDL output

- /home/lucy/netcdf/netcdf-c-4.7.4/ncdump/ncdump simple_xy.nc > test.cdl
- more test.cdl
- /home/lucy/netcdf/netcdf-c-4.7.4/ncgen/ncgen -b test.cdl
- ls
- cmp test.nc simple_xy.nc

Creating the C File

Create a C file using ncgen and the CDL output

- `/home/lucy/netcdf/netcdf-c-4.7.4/ncgen/ncgen -lc test.cdl > test.c`
- `more test.c`

► Start all over again!!!

Starting all over again!


- `gcc -I/home/lucy/local/include test.c -o test -L/home/lucy/local/lib -lnetcdf`
- `mv test.nc test2.nc`
- `./test`
- `ls`
- `cmp test.nc test2.nc`

File Reference: simple_xy_rd.c

This is a simple example which reads a small dummy array, which was written by simple_xy_wr.c. It is intended to illustrate the use of the netCDF C API.

- https://www.unidata.ucar.edu/software/netcdf/docs/simple__xy__rd_8c.html
- https://www.unidata.ucar.edu/software/netcdf/docs/simple__xy__rd_8c_source.html

Dependency graph for simple_xy_wr:



Reading the file simple_xy.nc

Create (copy!), compile and run the file simple_xy_rd.c

■ Compile the file simple_xy_rd.c



```
gcc -I/home/lucy/local/include simple_xy_rd.c -o simple_xy_rd -L/home/lucy/local/
```

■ Run the file simple_xy_rd



```
./simple_xy_rd
```



```
*** SUCCESS reading example file simple_xy.nc!
```

ncview



- Installation is fine.
- What does it do?

Building NetCDF
oooooo

File simple_xy_wr.c
oooooooooooooooooooo

File simple_xy_nc4_wr.c
●

File simple_nc4_wr.c
○

File sfc_pres_temp_wr.c
○

File pres_temp_4D_wr.c
○○



Building NetCDF
oooooo

File simple_xy_wr.c
oooooooooooooooooooo

File simple_xy_nc4_wr.c
o

File simple_nc4_wr.c
●

File sfc_pres_temp_wr.c
o

File pres_temp_4D_wr.c
oo



Building NetCDF
oooooo

File simple_xy_wr.c
oooooooooooooooooooo

File simple_xy_nc4_wr.c
o

File simple_nc4_wr.c
o

File sfc_pres_temp_wr.c
●

File pres_temp_4D_wr.c
oo



Building NetCDF
oooooo

File simple_xy_wr.c
oooooooooooooooooooo

File simple_xy_nc4_wr.c
o

File simple_nc4_wr.c
o

File sfc_pres_temp_wr.c
o

File pres_temp_4D_wr.c
●o



The ESiWACE1/2 projects have received funding from the European Union's Horizon 2020 research and innovation programme under grant agreement No **675191** and No **823988**



Disclaimer: This material reflects only the author's view and the EU-Commission is not responsible for any use that may be made of the information it contains