

2020 Summer School on Effective HPC for Climate and Weather

Input/Output and Middleware

Luciana Pedro, Julian Kunkel

Department of Computer Science, University of Reading

18 June 2020



1 NetCDF Files and C

2 NetCDF Utilities

3 Parallel I/O

4 Practising

Disclaimer: This material reflects only the author's view and the EU-Commission is not responsible for any use that may be made of the information it contains

Learning Objectives

- Execute programs in C that read and write NetCDF files in a metadata-aware manner
- Analyze, manipulate and visualise NetCDF data
- Implement an application that utilizes parallel I/O to store and analyze data

References


- The files and data used in this presentation were collected on the Unidata website.
 - ▶ <https://www.unidata.ucar.edu/>
- All files used here are available in the following Git Repository:
 - ▶ <https://github.com/ESiWACE/io-training>
- These files are also available with the NetCDF main installation, in the directory `examples`.
- For more information about how to install NetCDF in your personal computer, from scratch, check Section 5.

File Reference: `simple_xy_wr.c`

- This is an example program demonstrating a simple 2D write. It is intended to illustrate the use of the netCDF C API.

- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple__xy__wr_8c.html
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple__xy__wr_8c_source.html

- Dependency graph for `simple_xy_wr`:



File simple_xy_wr.c: Header and Constants Declaration

```
#include <stdlib.h>
#include <stdio.h>
#include <netcdf.h>

/* This is the name of the data file we will create. */
#define FILE_NAME "simple_xy.nc"

/* We are writing 2D data, a 6 x 12 grid. */
#define NDIMS 2
#define NX 6
#define NY 12

/* Handle errors by printing an error message and exiting with a
 * non-zero status. */
#define ERRCODE 2
#define ERR(e) {printf("Error: %s\n", nc_strerror(e)); exit(ERRCODE);}

int main()
{
    ...
    ...
    ...
}
```

File simple_xy_wr.c: Variables Declaration

```
...  
...  
...  
  
int main()  
{  
    /* When we create netCDF variables and dimensions, we get back an  
     * ID for each one. */  
    int ncid, x_dimid, y_dimid, varid;  
    int dimids[NDIMS];  
  
    /* This is the data array we will write. It will be filled with a  
     * progression of numbers for this example. */  
    int data_out[NX][NY];  
  
    /* Loop indexes, and error handling. */  
    int x, y, retval;  
  
    ...  
    ...  
    ...  
}
```

File simple_xy_wr.c: Creating (loading!) Data

```
...

int main()
{
    ...

    /* Create some pretend data. If this wasn't an example program, we
     * would have some real data to write, for example, model
     * output. */
    for (x = 0; x < NX; x++)
        for (y = 0; y < NY; y++)
            data_out[x][y] = x * NY + y;

    ...
}
```


File simple_xy_wr.c: Creating the NetCDF file

```
...

int main()
{
    ...

    /* Always check the return code of every netCDF function call. In
     * this example program, any retval which is not equal to NC_NOERR
     * (0) will cause the program to print an error message and exit
     * with a non-zero return code. */

    /* Create the file. The NC_CLOBBER parameter tells netCDF to
     * overwrite this file, if it already exists.*/
    if ((retval = nc_create(FILE_NAME, NC_CLOBBER, &ncid)))
        ERR(retval);

    ...
}
```

File simple_xy_wr.c: Defining the Dimensions

```
...

int main()
{
    ...

    /* Define the dimensions. NetCDF will hand back an ID for each. */
    if ((retval = nc_def_dim(ncid, "x", NX, &x_dimid)))
        ERR(retval);
    if ((retval = nc_def_dim(ncid, "y", NY, &y_dimid)))
        ERR(retval);

    /* The dimids array is used to pass the IDs of the dimensions of
       * the variable. */
    dimids[0] = x_dimid;
    dimids[1] = y_dimid;

    ...
}
```

File simple_xy_wr.c: Defining a Variable

```

...

int main()
{
    ...

    /* Define the variable. The type of the variable in this case is
     * NC_INT (4-byte integer). */
    if ((retval = nc_def_var(ncid, "data", NC_INT, NDIMS,
                            dimids, &varid)))
        ERR(retval);

    /* End define mode. This tells netCDF we are done defining
     * metadata. */
    if ((retval = nc_enddef(ncid)))
        ERR(retval);

    ...
}

```

File simple_xy_wr.c: Writing Data into the File

```

...

int main()
{
    ...

    /* Write the pretend data to the file. Although netCDF supports
     * reading and writing subsets of data, in this case we write all
     * the data in one operation. */
    if ((retval = nc_put_var_int(ncid, varid, &data_out[0][0])))
        ERR(retval);

    /* Close the file. This frees up any internal netCDF resources
     * associated with the file, and flushes any buffers. */
    if ((retval = nc_close(ncid)))
        ERR(retval);

    ...
}

```

File simple_xy_wr.c: Getting SUCCESS!



```
...  
  
int main()  
{  
    ...  
  
    printf("*** SUCCESS writing example file simple_xy.nc!\n");  
    return 0;  
}
```

Compiling and Running the File `simple_xy_wr.c`

- Create (copy!) and compile the file `simple_xy_wr.c`.

- ▶ `gcc simple_xy_wr.c -o simple_xy_wr $(nc-config --libs --cflags)`

- Run the file `simple_xy_wr`.

- ▶ `./simple_xy_wr`

- ▶ `*** SUCCESS writing example file simple_xy.nc!`


- Check that the file `simple_xy.nc` is in your directory.

File Reference: `simple_xy_rd.c`

- This is a simple example which reads a small dummy array that was written by `simple_xy_wr.c`.

- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple__xy__rd_8c.html
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple__xy__rd_8c_source.html

- Dependency graph for `simple_xy_rd`:



File simple_xy_rd.c

```
int main()
{
    /* Open the file. NC_NOWRITE tells netCDF we want read-only access
     * to the file.*/
    if ((retval = nc_open(FILE_NAME, NC_NOWRITE, &ncid)))
        ERR(retval);

    /* Get the varid of the data variable, based on its name. */
    if ((retval = nc_inq_varid(ncid, "data", &varid)))
        ERR(retval);

    /* Read the data. */
    if ((retval = nc_get_var_int(ncid, varid, &data_in[0][0])))
        ERR(retval);

    /* Check the data. */
    for (x = 0; x < NX; x++)
        for (y = 0; y < NY; y++)
            if (data_in[x][y] != x * NY + y)
                return ERRCODE;


    /* Close the file, freeing all resources. */
    if ((retval = nc_close(ncid)))
        ERR(retval);
}
```


Reading the File `simple_xy.nc`

- Check that the file `simple_xy.nc` is in your directory.
- Create (copy!), compile and run the file `simple_xy_rd.c`.
 - ▶ `gcc simple_xy_rd.c -o simple_xy_rd $(nc-config --libs --cflags)`
- Run the file `simple_xy_rd`.
 - ▶ `./simple_xy_rd`
 - ▶ `*** SUCCESS reading example file simple_xy.nc!`

ncdump and ncgen


- ncdump and ncgen are inverses:



- Used together, ncdump and ncgen can accomplish simple netCDF manipulations with little or no programming.

Editing a NetCDF File

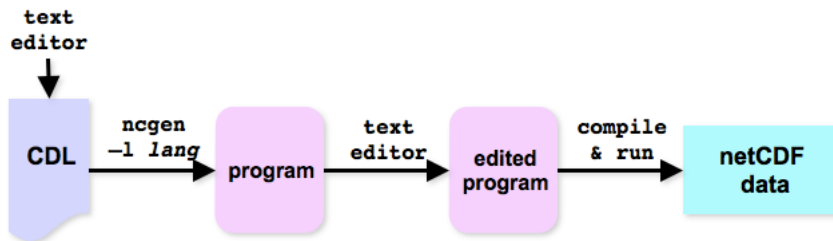
- To edit metadata or data in a netCDF file:



- ▶ Use `ncdump` to convert netCDF file to CDL.
- ▶ Use a text editor to make desired change to CDL.
- ▶ Use `ncgen` to turn modified CDL back into netCDF file.
- ▶ **Note:** This option is not practical for huge netCDF files or if one intend to modify lots of files. For that, need to write a program using netCDF library.

Creating a NetCDF File

- To create a new netCDF file with lots of metadata:



- ▶ Use a text editor to write a CDL file with lots of metadata but little or no data.
- ▶ Use `ncgen` to generate corresponding C or Fortran program for writing netCDF.
- ▶ Insert appropriate netCDF **var_put** calls for writing data.
- ▶ Compile and run program to create netCDF file.
- ▶ Use `ncdump` to verify result.

Using ncdump

- Inspect the file `simple_xy.nc` using `ncdump`

- ▶ `ncdump simple_xy.nc`

NetCDF CDL Format

```
netcdf simple_xy {  
  dimensions:  
    x = 6 ;  
    y = 12 ;  
  variables:  
    int data(x, y) ;  
  data:  
  
    data =  
      0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11,  
      12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23,  
      24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35,  
      36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47,  
      48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59,  
      60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71 ;  
}
```

Using ncgen

■ Create a NetCDF file using ncgen and the CDL output

▶ `ncdump simple_xy.nc > simple_xy_test.cdl`

▶ `more simple_xy_test.cdl`

▶ `ncgen -b simple_xy_test.cdl`

▶ `cmp simple_xy_test.nc simple_xy.nc`

Creating the C File

■ Create a C file using ncgen and the CDL output


- ▶ `ncgen -lc simple_xy_test.cdl > simple_xy_test.c`
- ▶ `more simple_xy_test.c`
- ▶ What is the difference between the files `simple_xy_test.c` and `simple_xy_wr.c`?
 - ▶ `cmp simple_xy_test.c simple_xy_wr.c`
 - ▶ `meld simple_xy_test.c simple_xy_wr.c`

Starting All Over Again!

- `gcc simple_xy_test.c $(nc-config --libs --cflags)`
`-o simple_xy_test -L/home/username/local/lib -lnetcdf`
- `mv simple_xy_test.nc simple_xy_test2.nc`
- `./simple_xy_test`
- `cmp simple_xy_test.nc simple_xy_test2.nc`

Using ncview


■ `ncview simple_xy.nc`




```
netcdf simple_xy {  
  dimensions:  
    x = 6 ;  
    y = 12 ;  
  variables:  
    int data(x, y) ;  
  data:  
  
    data =  
      0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11,  
      12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23,  
      24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35,  
      36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47,  
      48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59,  
      60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71 ;  
}
```

Global Potential Vegetation Dataset

- File `vegtype_5min.nc` – NetCDF 5 min data
- File `vegtype_0.5.nc` – NetCDF data aggregated to a 0.5 deg resolution



File `vegtype_5min.nc`



File `vegtype_0.5.nc`

Files available at <http://nelson.wisc.edu/sage/data-and-models/global-potential-vegetation/index.php>

Northern Hemisphere EASE-Grid Weekly Snow Cover

■ File `snowcover.mon.mean.nc`



January



February




March




April



May



June




July




August



September



October



November



December

File available at <https://ps1.noaa.gov/data/gridded/data.snowcover.html>.

Parallel I/O Example – MPI-IO

```
#include <stdio.h>
#include <mpi.h>
int main(int argc, char *argv[])
{
    MPI_File fh;
    int buf[1000], rank;
    MPI_Init(0,0);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_File_open(MPI_COMM_WORLD, "test.out",
        MPI_MODE_CREATE|MPI_MODE_WRONLY,
        MPI_INFO_NULL, &fh);
    if (rank == 0)
        MPI_File_write(fh, buf, 1000, MPI_INT, MPI_STATUS_IGNORE);
    MPI_File_close(&fh);
    MPI_Finalize();
    return 0;
}
```

Parallel I/O Example – NetCDF4 – Part I

```

/* Initialize MPI. */
MPI_Init(&argc,&argv);
MPI_Comm_size(MPI_COMM_WORLD, &mpi_size);
MPI_Comm_rank(MPI_COMM_WORLD, &mpi_rank);
MPI_Get_processor_name(mpi_name, &mpi_namelen);

if (mpi_rank == 1)
    printf("\n*** tst_parallel testing very basic parallel access.\n");

/* Create a parallel netcdf-4 file. */
if ((res = nc_create_par(FILE, NC_NETCDF4|NC_MPIIO, comm,
    info, &ncid))) ERR;

/* Create two dimensions. */
if ((res = nc_def_dim(ncid, "d1", DIMSIZE, dimids))) ERR;
if ((res = nc_def_dim(ncid, "d2", DIMSIZE, &dimids[1]))) ERR;

/* Create one var. */
if ((res = nc_def_var(ncid, "v1", NC_INT, NDIMS, dimids, &v1id))) ERR;

if ((res = nc_enddef(ncid))) ERR;

```

Parallel I/O Example – NetCDF4 – Part II

```
/* Set up slab for this process. */
start[0] = mpi_rank * DIMSIZE/mpi_size;
start[1] = 0;
count[0] = DIMSIZE/mpi_size;
count[1] = DIMSIZE;

/* Create phoney data. We're going to write a 24x24 array of ints,
   in 4 sets of 144. */
for (i=mpi_rank*QTR_DATA; i < (mpi_rank+1)*QTR_DATA; i++)
    data[i] = mpi_rank;

/*if ((res = nc_var_par_access(ncid, vlid, NC_COLLECTIVE)))
    ERR;*/
if ((res = nc_var_par_access(ncid, vlid, NC_INDEPENDENT))) ERR;

/* Write slabs of phoney data. */
if ((res = nc_put_vara_int(ncid, vlid, start, count,
    &data[mpi_rank*QTR_DATA]))) ERR;

/* Close the netcdf file. */
if ((res = nc_close(ncid))) ERR;

/* Shut down MPI. */
MPI_Finalize();

return 0;
```

Files for Practising

■ File simple_xy_nc4

- ▶ Write/Read the simple_xy file with some of the features of netCDF-4.
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_nc4_wr_8c.html
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_nc4_rd_8c.html

■ File simple_nc4

- ▶ Write/Read a file demonstrating some of the features of netCDF-4.
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_nc4_wr_8c.html
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/simple_nc4_rd_8c.html

■ File sfc_pres_temp_wr

- ▶ This is an example program which writes/reads surface pressure and temperatures.
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/sfc_pres_temp_wr_8c.html
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/sfc_pres_temp_rd_8c.html

■ File pres_temp_4D_wr

- ▶ This is an example program which writes/reads 4D pressure and temperatures.
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/pres_temp_4D_wr_8c.html
- ▶ https://www.unidata.ucar.edu/software/netcdf/docs/pres_temp_4D_rd_8c.html

Summary of Actions



- Inspect the read and write files in C code.
- Compile and run the write/read C files.
- Inspect the output NetCDF file (.nc) using `ncdump`.
- Create a CDL file for the NetCDF file.
- Recreate the NetCDF file using `ncgen` and the CDL file.
- Recreate the C file using `ncgen` and the CDL file.
- Visualize the data in the NetCDF file with `ncview`.

Appendix

Building NetCDF from Scratch

- The usual way of building netCDF requires the HDF5, zlib, curl and m4 libraries.
- Files for the libraries can be found in:

ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-4

- The following slides presents the steps for installing NetCDF in Ubuntu 18.04 and 20.04 for a user named **username**. Adapt the path to your own user.

Installing curl and m4

- apt-get install libcurl4-openssl-dev
- apt-get install m4

Installing zlib

- `wget ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-4/zlib-1.2.8.tar.gz`
 - ▶ Newest version to later use `ncview`
 - ▶ `wget https://sourceforge.net/projects/libpng/files/zlib/1.2.9/zlib-1.2.9.tar.gz`
- `tar -xvzf zlib-1.2.8.tar.gz`
- `cd zlib-1.2.8`
- `mkdir /home/username/local/`
- `./configure --prefix=/home/username/local/`
- `make check install`

Installing HDF5

- `wget ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-4/hdf5-1.8.13.tar.gz`
- `tar -xvzf hdf5-1.8.13.tar.gz`
- `cd hdf5-1.8.13`
- `./configure --with-zlib=/home/username/local/ --prefix=/home/username/local/`
- `make`
- `make check`
- `make install`
 - ▶ `make check install`
 - ▶ If not done separately, it might not work!

Installing NetCDF

- Check the latest version at
`https://www.unidata.ucar.edu/downloads/netcdf/`
- `wget ftp://ftp.unidata.ucar.edu/pub/netcdf/netcdf-c-4.7.4.tar.gz`
- `tar -xvzf netcdf-c-4.7.4.tar.gz`
- `cd netcdf-c-4.7.4`
- `CPPFLAGS=-I/home/username/local/include`
`LDLFLAGS=-L/home/username/local/lib ./configure`
`--prefix=/home/username/local`
- `make check install`

Finishing the Set Up

■ Link the NetCDF library

- ▶ `export LD_LIBRARY_PATH=/home/username/local/lib/`
- ▶ `sudo ldconfig`

■ Create a new directory (for instance, `/home/username/example`) and create the file from the given source using an editor of your choice.


File Reference: `simple_xy_wr.c`

This is an example program demonstrating a simple 2D write. It is intended to illustrate the use of the netCDF C API.

■ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_wr_8c.html

■ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_wr_8c_source.html

Dependency graph for `simple_xy_wr`:



Compiling and running the file `simple_xy_wr.c`

- Create (copy!) and compile the file `simple_xy_wr.c`.

- ▶ `gcc simple_xy_wr.c -o simple_xy_wr $(nc-config --libs --cflags)`

- Run the file `simple_xy_wr`.

- ▶ `./simple_xy_wr`

- ▶ `*** SUCCESS writing example file simple_xy.nc!`

- Check that the file `cmp test.nc simple_xy.nc` is in your directory.

Using ncdump

Inspect the output file `simple_xy.nc` using `ncdump`

```
■ ncdump simple_xy.nc
```

NetCDF CDL Format

```
netcdf simple_xy {  
  dimensions:  
    x = 6 ;  
    y = 12 ;  
  variables:  
    int data(x, y) ;  
  data:  
  
    data =  
      0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11,  
      12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23,  
      24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35,  
      36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47,  
      48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59,  
      60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71 ;  
}
```

Using ncgen

■ Create a NetCDF file using ncgen and the CDL output

- ▶ `ncdump simple_xy.nc > test.cdl`
- ▶ `more test.cdl`
- ▶ `ncgen -b test.cdl`
- ▶ `cmp test.nc simple_xy.nc`

Creating the C File

■ Create a C file using ncgen and the CDL output

- ▶ `ncgen -lc simple_xy_test.cdl > simple_xy_test.c`
- ▶ `more simple_xy_test.c`
- ▶ `cmp simple_xy_test.c simple_xy_wr.c`

Starting all over again!

- gcc simple_xy_test.c -o simple_xy_test \$(nc-config --libs --cflags)
- mv simple_xy_test.nc simple_xy_test2.nc
- ./simple_xy_test
- cmp simple_xy_test.nc simple_xy_test2.nc

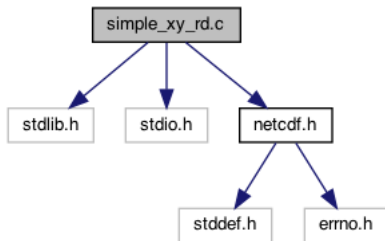
File Reference: `simple_xy_rd.c`

This is a simple example which reads a small dummy array, which was written by `simple_xy_wr.c`. It is intended to illustrate the use of the netCDF C API.

■ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_rd_8c.html

■ https://www.unidata.ucar.edu/software/netcdf/docs/simple_xy_rd_8c_source.html

Dependency graph for `simple_xy_wr`:



Reading the file `simple_xy.nc`

■ Check that the file `simple_xy.nc` is in your directory.

■ Create (copy!), compile and run the file `simple_xy_rd.c`

```
▶ gcc simple_xy_rd.c -o simple_xy_rd $(nc-config --libs --cflags)
```

■ Run the file `simple_xy_rd`

```
▶ ./simple_xy_rd
```

```
▶ *** SUCCESS reading example file simple_xy.nc!
```

The ESiWACE1/2 projects have received funding from the European Union's Horizon 2020 research and innovation programme under grant agreement No **675191** and No **823988**



Disclaimer: This material reflects only the author's view and the EU-Commission is not responsible for any use that may be made of the information it contains